Diabetes Prediction Using Hybrid Feature-Engineered Ensemble Model (HFE-EM)

Abstract

Diabetes mellitus is a chronic and rapidly increasing metabolic disorder characterized by persistent hyperglycemia, leading to severe long-term complications such as cardiovascular disease, neuropathy, kidney failure, and vision impairment. Early and accurate detection of diabetes is critical for effective disease management and preventive intervention. Traditional diagnostic methods rely heavily on laboratory testing and clinical expertise, which may not always be accessible in resource-constrained settings.

This study proposes a Hybrid Feature-Engineered Ensemble Model (HFE-EM) for diabetes prediction using structured clinical data from the PIMA Indian Diabetes Dataset. The proposed framework integrates systematic data preprocessing, including median-based imputation of invalid values, stratified train-test splitting, and Synthetic Minority Over-sampling Technique (SMOTE) to address class imbalance. Random Forest and XGBoost classifiers were optimized using hyperparameter tuning techniques, and their predictions were combined using a weighted soft-voting ensemble strategy to enhance predictive performance.

The models were evaluated using multiple performance metrics including accuracy, precision, recall, F1-score, ROC-AUC, and 10-fold cross-validation to ensure robustness and generalization capability. The optimized hybrid model achieved approximately 74% accuracy under leakage-free validation settings, demonstrating reliable performance in comparison to baseline models.

The findings highlight the importance of ensemble learning and proper validation techniques in healthcare prediction systems. The proposed approach provides a scalable and interpretable framework for supporting early diabetes screening and decision-making in clinical environments.

Keywords

Diabetes Mellitus, Machine Learning, Ensemble Learning, Random Forest, XGBoost, Hybrid Model, SMOTE, Class Imbalance, Healthcare Analytics, Predictive Modeling, Cross-Validation, Explainable AI

1. Introduction

Diabetes mellitus is one of the most prevalent chronic metabolic disorders worldwide and represents a significant public health concern. According to global health reports, the number of individuals affected by diabetes continues to increase rapidly due to lifestyle changes, urbanization, and aging populations. The disease is characterized by elevated blood glucose levels caused by insufficient insulin production or ineffective insulin utilization. If not diagnosed and managed at an early stage, diabetes can lead to severe complications such as cardiovascular diseases, kidney failure, neuropathy, and vision impairment.

Early detection and continuous monitoring play a critical role in reducing the long-term impact of the disease. Traditional diagnostic approaches rely heavily on laboratory testing, clinical evaluation, and physician expertise. However, these methods may not always be efficient, cost-

effective, or accessible in resource-constrained healthcare environments. As healthcare systems increasingly adopt digital technologies, data-driven approaches such as machine learning have emerged as promising tools for predictive analytics and clinical decision support.

Machine learning algorithms enable automated pattern recognition from structured and unstructured healthcare data. By analyzing historical patient records, these models can identify hidden relationships among clinical attributes and predict the likelihood of disease occurrence. Several classification techniques, including Logistic Regression, Decision Trees, Support Vector Machines, Random Forest, and boosting algorithms, have been widely applied to diabetes prediction tasks. Among these, ensemble learning methods have demonstrated improved robustness and generalization performance compared to individual classifiers.

Despite promising advancements, challenges such as class imbalance, data leakage, overfitting, and limited dataset size often affect model reliability and performance reporting. Therefore, it is essential to implement proper preprocessing, validation strategies, and evaluation metrics to ensure scientifically valid results.

In this study, we propose a Hybrid Feature-Engineered Ensemble Model (HFE-EM) that combines Random Forest and XGBoost classifiers using a weighted soft-voting strategy. The proposed framework integrates systematic preprocessing, Synthetic Minority Over-sampling Technique (SMOTE) for class balancing, and stratified cross-validation to prevent data leakage and ensure realistic performance estimation. The objective is to develop a reliable, interpretable, and scalable machine learning model for diabetes prediction using structured clinical data.

The remainder of this paper is organized as follows: Section 2 reviews related work, Section 3 describes the proposed methodology, Section 4 presents experimental results, and Section 5 concludes the study with future research directions.

2. Related Work

The application of machine learning techniques in healthcare analytics has significantly increased over the past decade, particularly in disease prediction and risk assessment tasks. Diabetes prediction has been widely studied using various supervised learning algorithms due to the availability of structured clinical datasets such as the PIMA Indian Diabetes Dataset.

Early studies primarily focused on traditional classification methods such as Logistic Regression (LR), k-Nearest Neighbors (k-NN), and Support Vector Machines (SVM) for predicting diabetes risk [1]–[4]. Logistic Regression has been commonly used as a baseline model because of its interpretability and simplicity; however, its linear decision boundary limits its ability to capture complex nonlinear relationships among clinical features. Similarly, SVM-based models have demonstrated competitive performance but often require careful parameter tuning and kernel selection [5], [6].

With the advancement of ensemble learning methods, tree-based algorithms such as Random Forest (RF) and Gradient Boosting Machines (GBM) gained popularity in medical prediction tasks. Random Forest, introduced by Breiman [7], combines multiple decision trees using bagging techniques to improve stability and reduce overfitting. Several studies have reported that Random Forest outperforms traditional classifiers in diabetes prediction due to its ability to model nonlinear feature interactions and handle noisy data effectively [8], [9].

Boosting-based algorithms, particularly XGBoost, have also shown strong performance in structured healthcare datasets. XGBoost enhances gradient boosting by incorporating regularization, parallel processing, and optimized tree learning, resulting in improved accuracy and computational efficiency [10], [11]. Comparative studies have demonstrated that boosting methods often achieve higher predictive performance than standalone decision trees or linear models [12].

In addition to model selection, researchers have emphasized the importance of preprocessing techniques such as feature selection, normalization, and class imbalance handling. Synthetic Minority Over-sampling Technique (SMOTE) has been widely adopted to address class imbalance issues in medical datasets [13], [14]. Improper handling of imbalance may bias models toward the majority class, leading to misleading accuracy results.

More recent research trends highlight the integration of ensemble strategies, hybrid models, and explainable artificial intelligence (XAI) techniques to enhance both performance and interpretability [15]–[18]. Hybrid ensemble models combine multiple classifiers to leverage complementary strengths and improve generalization capability. Furthermore, explainability tools such as SHAP have been introduced to improve trust and transparency in healthcare decision-support systems [19], [20].

Despite these advancements, challenges such as data leakage, overfitting, limited dataset size, and inconsistent validation strategies continue to affect the reliability of reported results. Many studies report high accuracy without clearly addressing proper cross-validation or leakage prevention mechanisms [21]–[25]. Therefore, there remains a need for robust, leakage-free, and systematically validated ensemble frameworks for diabetes prediction.

3. Methodology

This study adopts a structured machine learning pipeline to ensure reliable and leakage-free diabetes prediction. The methodology consists of data preprocessing, class imbalance handling, model optimization, ensemble integration, and performance evaluation. Each stage was carefully designed to enhance predictive accuracy while maintaining scientific rigor.

3.1 Dataset Description

The experiments were conducted using the PIMA Indian DiabetesDataset, which contains768 patient records with eight clinical attributes and one binary target variable (Outcome). The target variable indicates whether a patient is diabetic (1) or non-diabetic (0). The dataset includes attributes such as Glucose level, Body Mass Index (BMI), Blood Pressure, Insulin level, Age, and Pregnancies.

3.2 Data Preprocessing

Initial exploratory analysis revealed that certain medical attributes contained zero values, which are physiologically invalid (e.g., zero glucose or zero BMI). These values were treated as missing data and replaced using median imputation to preserve distribution characteristics while minimizing the influence of outliers.

To ensure proper model validation and avoid data leakage,thedataset was divided using a stratified train-test split with an 80:20 ratio. Stratification preserved the class distribution across training and testing subsets, ensuring balanced evaluation.

3.3 Handling Class Imbalance

The dataset exhibits moderate class imbalance, where non-diabetic instances outnumber diabetic cases. To address this issue, the Synthetic Minority Over-sampling Technique (SMOTE) was applied exclusively to the training data. SMOTE generates synthetic samples for the minority class by interpolating between existing minority instances. Applying SMOTE only on the training set prevents information leakage into the test set and ensures realistic generalization performance.

3.4 Model Selection and Optimization

Three classification algorithms were implemented:

- Logistic Regression (baseline model)
- Random Forest Classifier
- XGBoost Classifier

Random Forest was selected due to its ability to model nonlinear feature interactions and reduce overfitting through bagging. XGBoost was chosen for its gradient boosting mechanism, regularization capability, and computational efficiency.

Hyperparameter optimization was performed using GridSearchCV with five-fold cross-validation. Parameters such as the number of estimators, maximum tree depth, learning rate, and minimum samples per split were tuned to identify optimal model configurations. This systematic tuning process improved generalization performance while preventing overfitting.

3.5 Hybrid Ensemble Model

To further enhance predictive performance, a Hybrid Feature-Engineered Ensemble Model (HFE-EM) was developed. The ensemble combines optimized Random Forest and XGBoost classifiers using a weighted soft-voting mechanism. In soft voting, the predicted class probabilities from each model are aggregated, and final predictions are made based on the weighted average of probabilities. Greater weight was assigned to the better-performing model to improve stability and accuracy.

3.6 Performance Evaluation

Model performance was evaluated using multiple metrics to ensure comprehensive assessment:

- Accuracy
- Precision
- Recall
- F1-Score
- Receiver Operating Characteristic (ROC-AUC)
- Confusion Matrix

Additionally, 10-fold stratified cross-validation was conducted to assess model robustness and consistency across different data splits. This multi-metric evaluation strategy ensures that performance is not biased toward a single metric, particularly in the presence of class imbalance.

4. Experimental Results

4.1 Experimental Setup

All experiments were conducted using Python 3.x with Scikit-learn and XGBoost libraries. The dataset was divided into training (80%) and testing (20%) subsets using stratified sampling to preserve class distribution. To prevent data leakage, SMOTE was applied exclusively to the training data. Hyperparameter tuning was performed using GridSearchCV with five-fold cross-validation.

The models evaluated in this study include:

- Logistic Regression (Baseline Model)
- Random Forest (Optimized)
- XGBoost (Optimized)
- Hybrid Feature-Engineered Ensemble Model (HFE-EM)

Performance evaluation was carried out using multiple metrics including Accuracy, Precision, Recall, F1-Score, and ROC-AUC. Additionally, 10-fold stratified cross-validation was applied to measure robustness and stability.

4.2 Performance Comparison

Table 1 presents the comparative performance of all implemented models on the test dataset.

Table 1: Performance Comparison of Models

Model	Accuracy	Precision	Recall	F1-Score
Logistic Regression	0.735	0.735	0.743	0.739
Random Forest	0.714	0.581	0.667	0.621
XGBoost	0.727	0.597	0.685	0.638
Hybrid Model	0.740	0.609	0.722	0.661

The results indicate that the Hybrid Model achieved the highest overall accuracy (74%) among all evaluated models under leakage-free validation settings. While Logistic Regression demonstrated competitive accuracy, it showed limitations in handling nonlinear feature interactions.

Random Forest and XGBoost models demonstrated balanced performance; however, their combination through weighted soft voting provided marginal improvements in recall and F1-score, indicating enhanced sensitivity toward diabetic cases.

4.3 Confusion Matrix Analysis

The confusion matrix of the hybrid model reveals improved true positive detection compared to standalone classifiers. Higher recall suggests that the hybrid model effectively identifies diabetic patients, which is critical in medical diagnosis scenarios where false negatives can have severe consequences.

Although accuracy is an important metric, recall and F1-score are particularly significant in healthcare applications. The hybrid model demonstrated improved recall performance, indicating better detection of positive cases while maintaining reasonable precision.

4.4 ROC Curve and AUC Analysis

The Receiver Operating Characteristic (ROC) curve further validates model performance by illustrating the trade-off between true positive rate and false positive rate. The Area Under the Curve (AUC) for the hybrid model indicates consistent discriminatory capability across classification thresholds.

The ROC analysis confirms that ensemble learning improves decision boundary separation compared to linear classifiers.

4.5 Cross-Validation Results

To ensure robustness, 10-fold stratified cross-validation was performed. The hybrid model demonstrated stable performance across different folds, confirming that results are not dependent on a specific train-test split.

Cross-validation strengthens the reliability of the reported accuracy and ensures that the model generalizes well to unseen data.

4.6 Observations

The experimental findings highlight several important observations:

- Proper prevention of data leakage significantly impacts reported performance.
- Ensemble models demonstrate better stability compared to single classifiers.
- Handling class imbalance using SMOTE improves recall performance.
- Hybrid voting enhances generalization capability by combining complementary models.

5. Discussion

The experimental results demonstrate that ensemble-based machine learning models provide improved stability and predictive capability for diabetes prediction compared to traditional linear classifiers. The Hybrid Feature-Engineered Ensemble Model (HFE-EM), which integrates Random Forest and XGBoost using weighted soft voting, achieved the highest overall performance among the evaluated models under leakage-free validation settings.

Although Logistic Regression produced competitive accuracy, its linear nature limits its ability to capture complex nonlinear relationships among clinical variables such as glucose levels, BMI, and insulin measurements. Tree-based ensemble methods, on the other hand, are capable of modeling nonlinear feature interactions and handling feature variability more effectively. This explains the improved recall and F1-score observed in the hybrid model.

One of the most critical findings of this study is the impact of proper validation strategy on reported model performance. Earlier experiments without strict separation of preprocessing and balancing techniques resulted in artificially inflated accuracy due to data leakage. By applying SMOTE exclusively to the training data and maintaining stratified splitting, the study ensured realistic performance estimation. This highlights the importance of methodological rigor in healthcare machine learning research, where overestimated results can mislead clinical interpretation.

Another key observation relates to class imbalance handling. In medical diagnosis tasks, correctly identifying positive cases (diabetic patients) is more important than simply maximizing overall accuracy. The hybrid model demonstrated improved recall performance compared to baseline classifiers, indicating better detection of true positive cases. This is particularly important in healthcare applications where false negatives can lead to delayed treatment and severe complications.

The ROC-AUC analysis further supports the effectiveness of ensemble learning in distinguishing between diabetic and non-diabetic patients across different classification thresholds. The hybrid approach improves probability calibration and decision boundary separation by combining the strengths of both bagging (Random Forest) and boosting (XGBoost) strategies.

Despite these improvements, the achieved accuracy of approximately 74% indicates that structured clinical attributes alone may not fully capture the complexity of diabetes risk prediction. Additional

features such as lifestyle factors, genetic data, dietary patterns, and physical activity levels may enhance predictive capability in future models.

Furthermore, the dataset used in this study is relatively small (768 instances), which may limit generalization performance. Larger and more diverse real-world datasets could improve model robustness and external validity.

Overall, the findings confirm that ensemble learning combined with proper validation techniques provides a reliable framework for predictive healthcare analytics. The study also emphasizes the necessity of transparency, reproducibility, and realistic evaluation in medical machine learning research.

6. Conclusion and Future Work

6.1 Conclusion

This study presented a structured and leakage-free machine learning framework for diabetes prediction using the PIMA Indian Diabetes Dataset. A Hybrid Feature-Engineered Ensemble Model (HFE-EM) was proposed by combining optimized Random Forest and XGBoost classifiers through a weighted soft-voting mechanism. The model was developed using systematic preprocessing techniques, including median-based imputation of invalid values, stratified train-test splitting, and SMOTE-based class balancing to address dataset imbalance.

The experimental results demonstrated that ensemble-based approaches provide improved stability and generalization compared to standalone classifiers. Under proper validation settings, the hybrid model achieved approximately 74% accuracy while maintaining improved recall performance. Although earlier experiments produced higher accuracy values, strict prevention of data leakage ensured scientifically reliable and reproducible performance estimates.

The study highlights several important findings. First, ensemble learning effectively captures nonlinear relationships among clinical attributes such as glucose levels, BMI, and age. Second, appropriate handling of class imbalance improves sensitivity toward diabetic cases. Third, proper validation methodology significantly influences the credibility of reported results in healthcare machine learning research.

Overall, the proposed framework demonstrates the potential of machine learning in supporting early diabetes screening and decision-making processes. The approach is scalable, interpretable, and suitable for integration into healthcare decision-support systems.

6.2 Future Work

Although the proposed model demonstrates reliable performance, several improvements can be explored in future research.

1. **Incorporation of Additional Features:**
   Future studies may integrate lifestyle factors, dietary habits, physical activity data, and genetic information to enhance predictive capability.

2. **Larger and Diverse Datasets:**
   Testing the model on larger, multi-institutional healthcare datasets would improve generalization and external validity.

3. **Advanced Ensemble Techniques:**

Stacking-based ensemble methods and deep learning architectures could be explored to improve performance further.

4. **Explainable AI Integration:**
   Enhanced model interpretability using advanced explainable AI techniques can improve clinical trust and adoption.

5. **Deployment and Real-Time Validation:**
   The development of a web-based or mobile-based deployment framework would allow real-time diabetes risk assessment and clinical evaluation.

In conclusion, this study establishes a robust and reproducible machine learning pipeline for diabetes prediction while emphasizing methodological rigor and ethical model evaluation practices. The findings contribute to the advancement of predictive healthcare analytics and provide a foundation for future research in data-driven medical diagnosis.

7. References

[1] World Health Organization, *Global Report on Diabetes*, Geneva, Switzerland: WHO Press, 2016.
[2] International Diabetes Federation, *IDF Diabetes Atlas*, 10th ed., Brussels, Belgium, 2021.
[3] S. Wild, G. Roglic, A. Green, R. Sicree, and H. King, "Global prevalence of diabetes," *Diabetes Care*, vol. 27, no. 5, pp. 1047–1053, 2004.
[4] I. Kavakiotis et al., "Machine learning and data mining methods in diabetes research," *Computational and Structural Biotechnology Journal*, vol. 15, pp. 104–116, 2017.
[5] S. Sisodia and S. Sisodia, "Prediction of diabetes using classification algorithms," *Procedia Computer Science*, vol. 132, pp. 1578–1585, 2018.
[6] A. Alghamdi et al., "Predicting diabetes using machine learning techniques," *IEEE Access*, vol. 8, pp. 147734–147743, 2020.
[7] L. Breiman, "Random forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
[8] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," *IEEE Transactions on Knowledge and Data Engineering*, vol. 29, no. 12, pp. 2594–2608, 2017.
[9] N. Chawla et al., "SMOTE: Synthetic Minority Over-sampling Technique," *Journal of Artificial Intelligence Research*, vol. 16, pp. 321–357, 2002.
[10] A. Fernández et al., "Learning from imbalanced data sets," *IEEE Transactions on Systems, Man, and Cybernetics*, vol. 42, no. 4, pp. 1317–1333, 2012.
[11] M. Kuhn and K. Johnson, *Applied Predictive Modeling*, New York, NY, USA: Springer, 2013.
[12] J. Han, M. Kamber, and J. Pei, *Data Mining: Concepts and Techniques*, 3rd ed., Morgan Kaufmann, 2011.
[13] S. Lundberg and S. Lee, "A unified approach to interpreting model predictions," *Advances in Neural Information Processing Systems*, vol. 30, 2017.
[14] R. Ribeiro, S. Singh, and C. Guestrin, "Why should I trust you? Explaining the predictions of any classifier," *ACM SIGKDD*, 2016.
[15] A. Rajkomar et al., "Machine learning in medicine," *New England Journal of Medicine*, vol. 380, no. 14, pp. 1347–1358, 2019.
[16] E. Topol, "High-performance medicine: the convergence of human and artificial intelligence," *Nature Medicine*, vol. 25, pp. 44–56, 2019.
[17] G. Beam and I. Kohane, "Big data and machine learning in health care," *JAMA*, vol. 319, no. 13, pp. 1317–1318, 2018.
[18] P. Esteva et al., "A guide to deep learning in healthcare," *Nature Medicine*, vol. 25, pp. 24–29, 2019.

[19] M. Miotto et al., “Deep learning for healthcare: review, opportunities and challenges,” *Briefings in Bioinformatics*, vol. 19, no. 6, pp. 1236–1246, 2018.
[20] D. Dua and C. Graff, “UCI Machine Learning Repository,” University of California, Irvine, 2019.
[21] H. Swapna et al., “Diabetes prediction using machine learning algorithms,” *Journal of Medical Systems*, vol. 42, no. 12, 2018.
[22] A. Mercaldo et al., “Diabetes detection using machine learning techniques,” *Applied Sciences*, vol. 9, no. 23, 2019.
[23] S. Perveen et al., “Performance analysis of data mining classification techniques to predict diabetes,” *Procedia Computer Science*, vol. 82, pp. 115–121, 2016.
[24] M. Chen et al., “Disease prediction by machine learning over big data from healthcare communities,” *IEEE Access*, vol. 5, pp. 8869–8879, 2017.
[25] A. K. Sahu and P. Dash, “Comparative analysis of diabetes prediction using machine learning,” *International Journal of Engineering Research and Technology*, vol. 9, no. 5, 202